

practice

November 19, 2024

```
[2]: import kagglehub

# Download latest version
path = kagglehub.dataset_download("vikrishnan/boston-house-prices")

print("Path to dataset files:", path)
```

Downloading from
https://www.kaggle.com/api/v1/datasets/download/vikrishnan/boston-house-prices?dataset_version_number=1...
100%|
12.8k/12.8k [00:00<00:00, 6.83MB/s]
Extracting files...
Path to dataset files: /Users/bruce/.cache/kagglehub/datasets/vikrishnan/boston-house-prices/versions/1

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns

import matplotlib.pyplot as plt
import plotly.express as px
import plotly.io as pio
pio.renderers.default = 'iframe'
from plotly import graph_objects as go
```

```
[2]: path = '/Users/bruce/.cache/kagglehub/datasets/vikrishnan/boston-house-prices/versions/1'
col_names = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV']
df = pd.read_csv(path + '/housing.csv', header=None, delimiter=r"\s+", names=col_names)
```

```
[3]: df
```

```
[3]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273.0	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273.0	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273.0	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273.0	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273.0	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2
..
501	21.0	391.99	9.67	22.4
502	21.0	396.90	9.08	20.6
503	21.0	396.90	5.64	23.9
504	21.0	393.45	6.48	22.0
505	21.0	396.90	7.88	11.9

[506 rows x 14 columns]

```
[4]: df.describe()
```

```
[4]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	AGE	DIS	RAD	TAX	PTRATIO	B	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000	

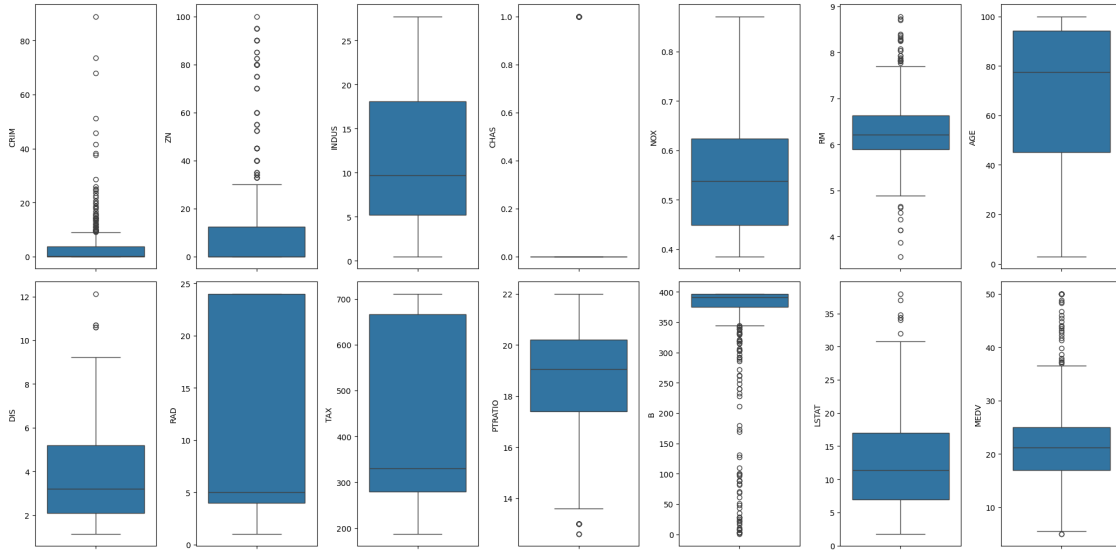
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT	MEDV
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[5]: df.isna().any(axis=0)
```


```
[5]: CRIM      False
      ZN       False
      INDUS   False
      CHAS    False
      NOX     False
      RM      False
      AGE     False
      DIS     False
      RAD     False
      TAX     False
      PTRATIO False
      B       False
      LSTAT   False
      MEDV    False
      dtype: bool
```

```
[6]: fig, axs = plt.subplots(figsize=(20,10), ncols=7, nrows=2)
      axs = axs.flatten()
      for i, col in enumerate(col_names):
          sns.boxplot(data=df, y=col, ax=axs[i])
      plt.tight_layout()
```




```
[253]: sns.displot(df['MEDV'], rug=True, kde=True)
```

```
[253]: <seaborn.axisgrid.FacetGrid at 0x2bc735870>
```




```
[255]: sns.displot(np.log(df['MEDV']) / df['RM'], rug=True, kde=True)
```

```
[255]: <seaborn.axisgrid.FacetGrid at 0x2bca1aa70>
```

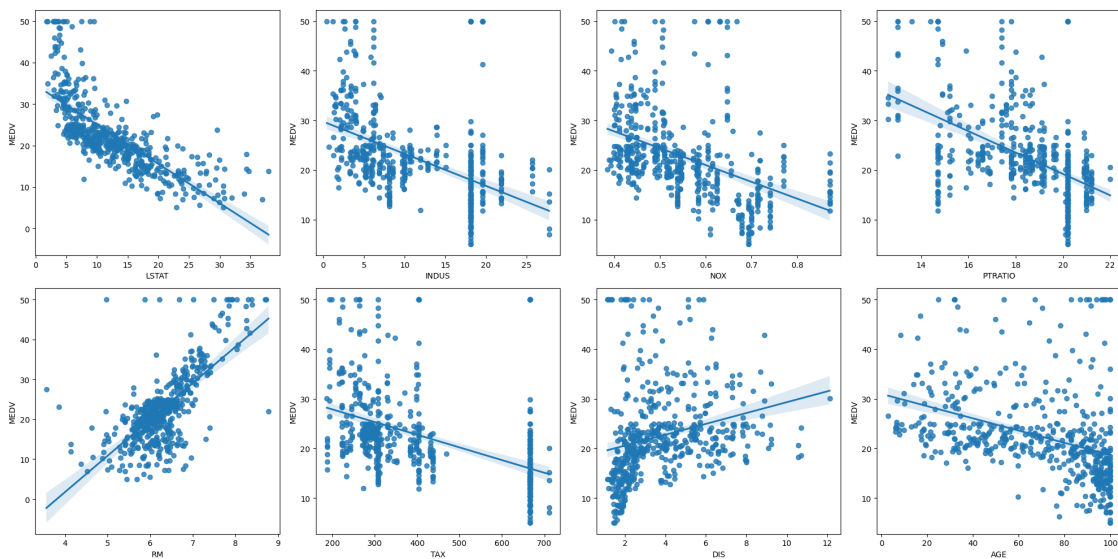


```
[258]: plt.figure(figsize=(20,10))
sns.heatmap(df.corr(), annot=True)
#px.imshow(df.corr(), text_auto=True)
```

[258]: <Axes: >



```
[260]: fig, axs = plt.subplots(ncols=4, nrows=2, figsize=(20, 10))
axs = axs.flatten()
for i, k in enumerate(['LSTAT', 'INDUS', 'NOX', 'PTRATIO', 'RM', 'TAX', 'DIS', 'AGE']):
    sns.regplot(y=df['MEDV'], x=df[k], ax=axs[i])
plt.tight_layout()
```



```
[9]: df_norm = (df - df.mean()) / df.std()
```

1 PCA


```
[10]: from sklearn import decomposition
```

```
pca = decomposition.PCA(n_components=10)
pca.fit(df_norm.drop(['MEDV'], axis=1))
```

```
[10]: PCA(n_components=10)
```


```
[17]: plt.figure(figsize=(20,10))
sns.heatmap(pd.DataFrame(pca.components_, columns=col_names[:-1]), annot=True)
```

```
[17]: <Axes: >
```



```
[13]: px.imshow(pca.components_, x=col_names[:-1], text_auto=True)
```

```
[342]: sns.lineplot(data=pca.explained_variance_, marker='o')
plt.show()
```

```
[232]: df_pca = pd.DataFrame(pca.transform(df_norm.drop(['MEDV'], axis=1)))
df_pca['MEDV'] = df['MEDV']
```

```
[233]: df_pca.corr()[['MEDV']]
```

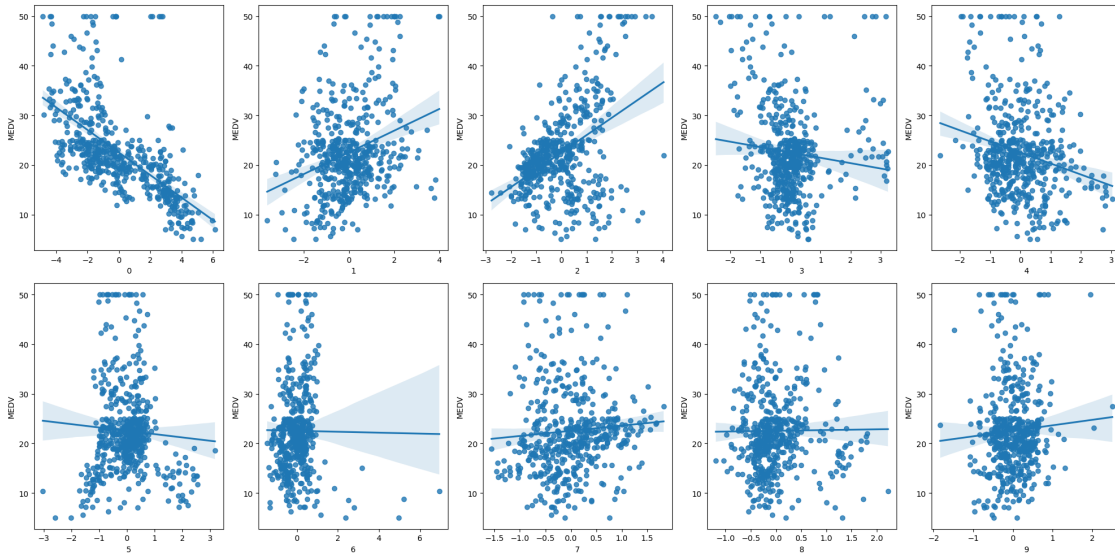
```
[233]:
```

	MEDV
0	-0.611745
1	0.285714
2	0.424334
3	-0.108814
4	-0.221845
5	-0.059122
6	-0.007503
7	0.071180
8	0.008551
9	0.056572
MEDV	1.000000

```
[261]: fig, axs = plt.subplots(nrows=2, ncols=5, figsize=(20,10))
axs = axs.flatten()

for i in range(10):
```

```
sns.regplot(y=df_pca['MEDV'], x=df_pca[i], ax=axes[i])
plt.tight_layout()
```



```
[17]: px.scatter(x=df['LSTAT'], y=df['MEDV'])
```

```
[18]: px.scatter(x=df['LSTAT'], y=X_pca[:,0])
```

2 OLS

```
[18]: from sklearn.model_selection import train_test_split
import statsmodels.api as sm
```

```
[19]: X = df[['LSTAT', 'DIS', 'TAX', 'PTRATIO', 'RM']]
y = df['MEDV']
X = sm.add_constant(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1,
↪shuffle=True)
```

```
[21]: import statsmodels.api as sm
```

```
[22]: model = sm.OLS(exog=X_train, endog=y_train)
results = model.fit() #_regularized(L1_wt=1, alpha=0) # L1_wt=1 means Lasso
```

```
[23]: results.summary()
```

```
[23]:
```

Dep. Variable:	MEDV	R-squared:	0.697
Model:	OLS	Adj. R-squared:	0.694
Method:	Least Squares	F-statistic:	206.5
Date:	Tue, 19 Nov 2024	Prob (F-statistic):	6.15e-114
Time:	22:14:49	Log-Likelihood:	-1378.6
No. Observations:	455	AIC:	2769.
Df Residuals:	449	BIC:	2794.
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	21.7375	4.261	5.102	0.000	13.364	30.111
LSTAT	-0.6095	0.051	-11.908	0.000	-0.710	-0.509
DIS	-0.7325	0.140	-5.239	0.000	-1.007	-0.458
TAX	-0.0054	0.002	-2.835	0.005	-0.009	-0.002
PTRATIO	-0.8121	0.127	-6.380	0.000	-1.062	-0.562
RM	4.5379	0.451	10.067	0.000	3.652	5.424

Omnibus:	188.784	Durbin-Watson:	1.954
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1030.092
Skew:	1.727	Prob(JB):	2.08e-224
Kurtosis:	9.512	Cond. No.	7.99e+03


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 7.99e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[34]: sns.scatterplot(x=y_train, y=results.resid)
```

```
[34]: <Axes: xlabel='MEDV', ylabel='None'>
```



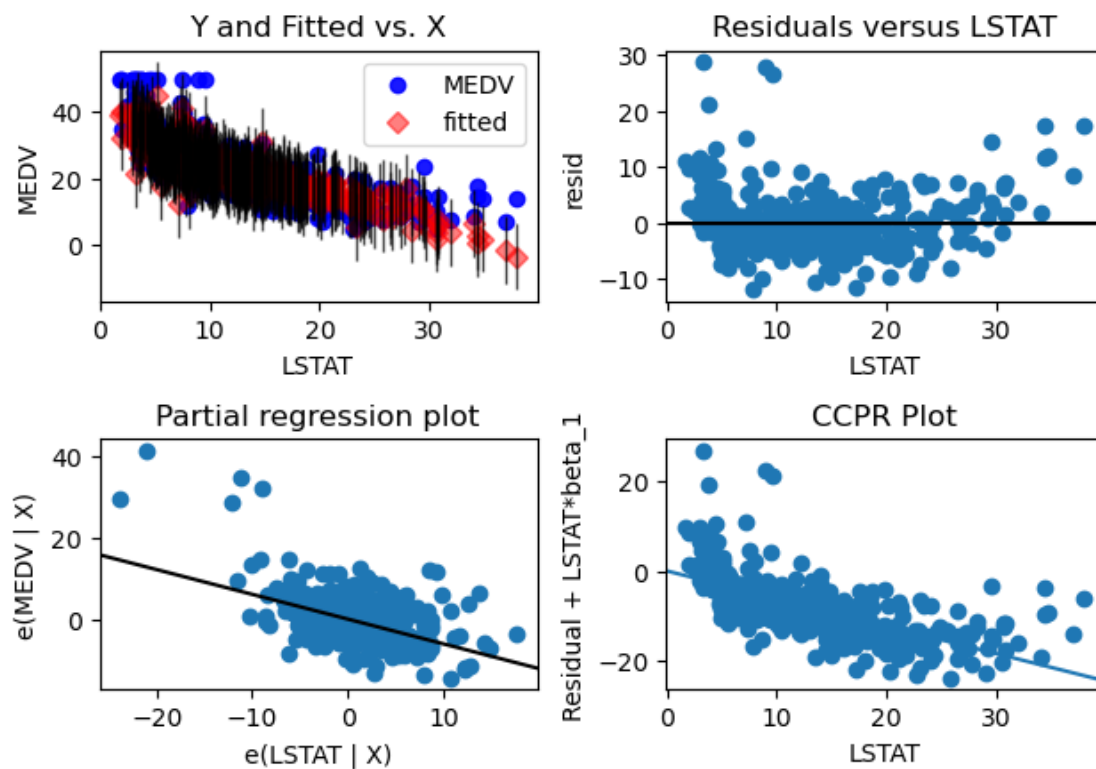
```
[24]: from sklearn.metrics import mean_squared_error
```

```
[25]: mean_squared_error(results.predict(X_test), y_test)
```

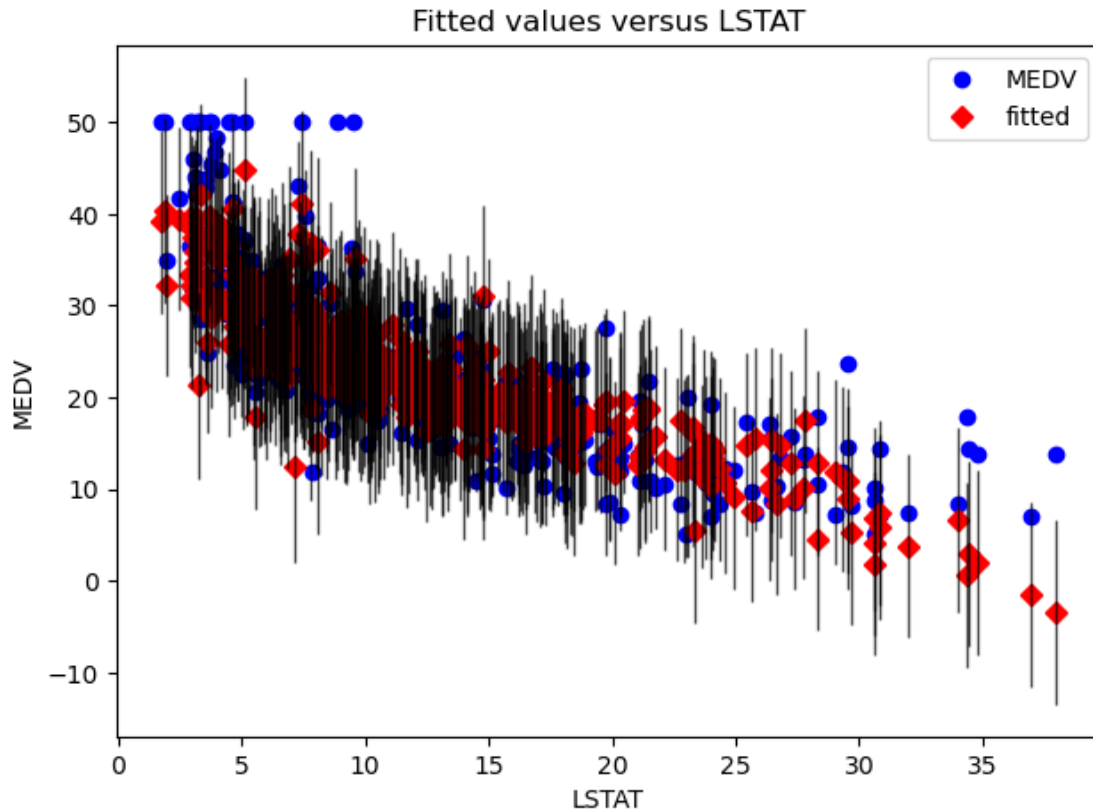
```
[25]: 31.31090608217717
```

```
[39]: fig = sm.graphics.plot_regress_exog(results, "LSTAT")  
fig.tight_layout(pad=1.0)
```

Regression Plots for LSTAT



```
[38]: fig = sm.graphics.plot_fit(results, "LSTAT")
fig.tight_layout(pad=1.0)
```



```
[42]: conf = results.get_prediction(X_test).summary_frame(alpha=0.05).
      ↪join(X_test['LSTAT'])
      conf.head(5)
```

```
[42]:
```

	mean	mean_se	mean_ci_lower	mean_ci_upper	obs_ci_lower	\
7	19.697067	0.701645	18.318151	21.075983	9.694338	
248	22.107377	0.556888	21.012946	23.201807	12.139881	
119	20.647399	0.378258	19.904023	21.390775	10.712320	
203	38.159159	0.650482	36.880792	39.437525	28.169794	
424	15.031669	0.492971	14.062852	16.000487	5.077183	


	obs_ci_upper	LSTAT
7	29.699796	19.15
248	32.074872	9.52
119	30.582478	13.61
203	48.148524	3.81
424	24.986155	17.16

```
[61]: x = np.linspace(0, 10, 100)
      y = np.sin(x)
      y_upper = y + 0.3
```

```
y_lower = y - 0.3

ax = sns.lineplot(x=x, y=y)
ax.fill_between(x, y_lower, y_upper, alpha=0.2, color='b', label='Confidence')
```


[61]: <matplotlib.collections.PolyCollection at 0x294e2f3d0>



```
[26]: res = y - results.predict(X)
```

```
[33]: plt.scatter(y, res)
```

[33]: <matplotlib.collections.PathCollection at 0x16a2caf80>




```
[28]: df['res'] = res
```

```
[29]: px.imshow(df.corr(), text_auto=True)
```

```
[30]: plt.scatter(df['RM'], res)
```

```
[30]: <matplotlib.collections.PathCollection at 0x16a032500>
```

3 GLM

see: <https://www.statsmodels.org/stable/glm.html>

```
[67]: log_model = sm.GLM(endog=y_train, exog=X_train, family=sm.families.  
      ↪Gaussian(link=sm.families.links.Log()))  
      log_results = log_model.fit()
```

```
[68]: log_results.summary()
```

```
[68]:
```

Dep. Variable:	MEDV	No. Observations:	455
Model:	GLM	Df Residuals:	449
Model Family:	Gaussian	Df Model:	5
Link Function:	Log	Scale:	18.406
Method:	IRLS	Log-Likelihood:	-1305.3
Date:	Tue, 19 Nov 2024	Deviance:	8264.3
Time:	22:44:53	Pearson chi2:	8.26e+03
No. Iterations:	8	Pseudo R-squ. (CS):	0.9701
Covariance Type:	nonrobust		


	coef	std err	z	P> z	[0.025	0.975]
const	3.0235	0.152	19.882	0.000	2.725	3.322
LSTAT	-0.0385	0.002	-15.845	0.000	-0.043	-0.034
DIS	-0.0299	0.005	-6.333	0.000	-0.039	-0.021
TAX	-0.0001	7.68e-05	-1.612	0.107	-0.000	2.67e-05
PTRATIO	-0.0242	0.004	-5.610	0.000	-0.033	-0.016
RM	0.1783	0.015	12.033	0.000	0.149	0.207

```
[69]: mean_squared_error(log_results.predict(X_test), y_test)
```

```
[69]: 28.42567629267149
```

```
[70]: plt.scatter(y_train, log_results.resid_response)
```

```
[70]: <matplotlib.collections.PathCollection at 0x2950bd5a0>
```



```
[100]: # --- Logit model ---
# sm.Logit(endog=, exog=)
```

```
[ ]:
```

4 RandomForest

```
[71]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
```

```
[72]: rf_cv = GridSearchCV(RandomForestRegressor(),
                          param_grid={'n_estimators': [100, 300, 500],
                                      'max_depth': [3, 5, 10],
                                      'max_features': [1.0, 'sqrt']},
                          scoring='neg_mean_squared_error',
                          cv=5,
                          verbose=0,
                          )
```

```
[282]: rf_cv.fit(X_train, y_train)
```

```
[282]: GridSearchCV(cv=5, estimator=RandomForestRegressor(),
                  param_grid={'max_depth': [3, 5, 10], 'max_features': [1.0, 'sqrt'],
                              'n_estimators': [100, 300, 500]},
                  scoring='neg_mean_squared_error')
```

```
[292]: #rf_cv.best_params_
pd.DataFrame(rf_cv.cv_results_).sort_values('rank_test_score')
```

```
[292]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
17	0.266116	0.055677	0.009936	0.000291	
15	0.045862	0.000494	0.002179	0.000047	
16	0.141573	0.004142	0.006410	0.000597	
14	0.336194	0.002969	0.009589	0.000132	
9	0.039489	0.000439	0.001889	0.000120	
13	0.208770	0.009322	0.006215	0.000291	
12	0.070385	0.001201	0.002541	0.000160	
6	0.054928	0.000717	0.002127	0.000042	
11	0.201186	0.004776	0.007659	0.000229	
8	0.264963	0.003692	0.007369	0.000204	
10	0.119924	0.002824	0.004712	0.000202	
7	0.157109	0.002981	0.004470	0.000101	
0	0.051108	0.012603	0.001796	0.000204	
1	0.138788	0.003875	0.004180	0.000202	

2	0.227359	0.003771	0.006639	0.000158
4	0.110966	0.003778	0.004272	0.000219
3	0.038823	0.000665	0.002038	0.000118
5	0.180174	0.003536	0.006502	0.000248

	param_max_depth	param_max_features	param_n_estimators	\
17	10	sqrt	500	
15	10	sqrt	100	
16	10	sqrt	300	
14	10	1.0	500	
9	5	sqrt	100	
13	10	1.0	300	
12	10	1.0	100	
6	5	1.0	100	
11	5	sqrt	500	
8	5	1.0	500	
10	5	sqrt	300	
7	5	1.0	300	
0	3	1.0	100	
1	3	1.0	300	
2	3	1.0	500	
4	3	sqrt	300	
3	3	sqrt	100	
5	3	sqrt	500	

	params	split0_test_score	\
17	{'max_depth': 10, 'max_features': 'sqrt', 'n_e...	-7.448951	
15	{'max_depth': 10, 'max_features': 'sqrt', 'n_e...	-7.070799	
16	{'max_depth': 10, 'max_features': 'sqrt', 'n_e...	-7.201673	
14	{'max_depth': 10, 'max_features': 1.0, 'n_esti...	-7.965726	
9	{'max_depth': 5, 'max_features': 'sqrt', 'n_es...	-8.194504	
13	{'max_depth': 10, 'max_features': 1.0, 'n_esti...	-7.966277	
12	{'max_depth': 10, 'max_features': 1.0, 'n_esti...	-7.491151	
6	{'max_depth': 5, 'max_features': 1.0, 'n_estim...	-8.616203	
11	{'max_depth': 5, 'max_features': 'sqrt', 'n_es...	-8.532719	
8	{'max_depth': 5, 'max_features': 1.0, 'n_estim...	-8.478953	
10	{'max_depth': 5, 'max_features': 'sqrt', 'n_es...	-8.721183	
7	{'max_depth': 5, 'max_features': 1.0, 'n_estim...	-8.421314	
0	{'max_depth': 3, 'max_features': 1.0, 'n_estim...	-12.063792	
1	{'max_depth': 3, 'max_features': 1.0, 'n_estim...	-11.976411	
2	{'max_depth': 3, 'max_features': 1.0, 'n_estim...	-12.149863	
4	{'max_depth': 3, 'max_features': 'sqrt', 'n_es...	-13.107641	
3	{'max_depth': 3, 'max_features': 'sqrt', 'n_es...	-12.410374	
5	{'max_depth': 3, 'max_features': 'sqrt', 'n_es...	-12.854066	

	split1_test_score	split2_test_score	split3_test_score	\
17	-12.689259	-13.078331	-13.908736	

15	-11.195586	-13.411535	-14.703290
16	-12.223230	-13.175793	-14.328282
14	-11.919455	-16.493886	-16.287242
9	-14.436129	-12.550843	-15.544291
13	-13.087440	-16.380269	-15.722388
12	-12.378430	-16.560149	-17.330716
6	-12.503808	-16.656268	-15.798471
11	-13.834109	-13.650729	-15.702470
8	-12.016245	-16.452318	-16.089019
10	-15.904770	-13.154462	-15.466326
7	-12.482940	-16.420491	-16.294311
0	-16.740463	-20.460051	-21.059465
1	-16.584222	-20.265343	-21.565937
2	-16.722707	-20.207380	-21.032545
4	-24.639471	-15.641116	-21.905448
3	-25.998506	-16.204044	-20.704805
5	-25.320643	-16.432822	-21.382617


	split4_test_score	mean_test_score	std_test_score	rank_test_score
17	-10.563946	-11.537845	2.323252	1
15	-11.781437	-11.632529	2.592948	2
16	-11.280187	-11.641833	2.439421	3
14	-10.736027	-12.680467	3.290585	4
9	-12.773048	-12.699763	2.507238	5
13	-10.448648	-12.721004	3.171871	6
12	-9.919546	-12.735999	3.776354	7
6	-10.380804	-12.791111	3.075895	8
11	-12.433978	-12.830801	2.389827	9
8	-11.313827	-12.870072	3.020826	10
10	-11.712183	-12.991785	2.627193	11
7	-11.450183	-13.013848	3.039241	12
0	-13.521545	-16.769063	3.597713	13
1	-13.514683	-16.781319	3.710199	14
2	-14.242345	-16.870968	3.396199	15
4	-19.132695	-18.885274	4.152085	16
3	-19.179704	-18.899486	4.541096	17
5	-19.587182	-19.115466	4.248963	18

```
[293]: mean_squared_error(rf_cv.predict(X_test), y_test)
```


```
[293]: 10.781680971439544
```

```
[294]: plt.scatter(y_test, y_test - rf.predict(X_test))
```

```
[294]: <matplotlib.collections.PathCollection at 0x2bf116320>
```



```
[300]: from sklearn.tree import plot_tree
plt.figure(figsize=(200,10))
plot_tree(rf_cv.best_estimator_.estimators_[0])
plt.show()
```



```
[147]: import xgboost
```

```
[172]: bst = xgboost.XGBRegressor(n_estimators=5000, max_depth=3, learning_rate=1e-3)
```

```
[173]: bst.fit(X_train, y_train)
```

```
[173]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                  colsample_bylevel=None, colsample_bynode=None,
                  colsample_bytree=None, device=None, early_stopping_rounds=None,
                  enable_categorical=False, eval_metric=None, feature_types=None,
                  gamma=None, grow_policy=None, importance_type=None,
                  interaction_constraints=None, learning_rate=0.001, max_bin=None,
```

```
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=3, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=5000, n_jobs=None,
num_parallel_tree=None, random_state=None, ...)
```

```
[174]: mean_squared_error(bst.predict(X_test), y_test)
```

```
[174]: 11.443298078843423
```

```
[177]: px.scatter(x=y_test, y=y_test - bst.predict(X_test))
```

5 OLS with formulas

```
[301]: import statsmodels.formula.api as smf
```

```
[321]: model_f = smf.ols(formula='np.log1p(MEDV) ~ LSTAT*RM + DIS + np.log1p(TAX) +  
↳ PTRATIO + 1', data=pd.concat([X_train, y_train], axis=1))
```

```
[322]: result_f = model_f.fit()
result_f.summary()
```

```
[322]:
```

Dep. Variable:	np.log1p(MEDV)	R-squared:	0.770
Model:	OLS	Adj. R-squared:	0.767
Method:	Least Squares	F-statistic:	250.0
Date:	Wed, 16 Oct 2024	Prob (F-statistic):	1.53e-139
Time:	23:05:23	Log-Likelihood:	114.88
No. Observations:	455	AIC:	-215.8
Df Residuals:	448	BIC:	-186.9
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	3.6643	0.256	14.315	0.000	3.161	4.167
LSTAT	0.0326	0.010	3.342	0.001	0.013	0.052
RM	0.2044	0.023	8.966	0.000	0.160	0.249
LSTAT:RM	-0.0114	0.002	-6.981	0.000	-0.015	-0.008
DIS	-0.0235	0.005	-4.437	0.000	-0.034	-0.013
np.log1p(TAX)	-0.1444	0.030	-4.865	0.000	-0.203	-0.086
PTRATIO	-0.0246	0.005	-5.092	0.000	-0.034	-0.015

Omnibus:	53.724	Durbin-Watson:	2.103
Prob(Omnibus):	0.000	Jarque-Bera (JB):	286.224
Skew:	0.310	Prob(JB):	7.04e-63
Kurtosis:	6.836	Cond. No.	2.57e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, $2.57e+03$. This might indicate that there are strong multicollinearity or other numerical problems.

```
[ ]:
```

6 More

```
[7]: # --- Flow ---  
# 1. Motivation; 2. Observation; 3. Conclusion
```

```
[79]: # --- k-fold CV for generalisation error ---  
from sklearn import datasets, linear_model  
from sklearn.model_selection import cross_val_score  
  
diabetes = datasets.load_diabetes()  
X = diabetes.data[:150]  
y = diabetes.target[:150]  
lasso = linear_model.Lasso()  
  
cross_val_score(lasso, X, y, cv=3, scoring='neg_mean_squared_error')
```

```
[79]: array([-3635.51152303, -3573.34242148, -6114.78229547])
```

```
[78]: import sklearn  
sklearn.metrics.SCORERS.keys()
```

```
[78]: dict_keys(['explained_variance', 'r2', 'max_error', 'neg_median_absolute_error',  
'neg_mean_absolute_error', 'neg_mean_absolute_percentage_error',  
'neg_mean_squared_error', 'neg_mean_squared_log_error',  
'neg_root_mean_squared_error', 'neg_mean_poisson_deviance',  
'neg_mean_gamma_deviance', 'accuracy', 'top_k_accuracy', 'roc_auc',  
'roc_auc_ovr', 'roc_auc_ovo', 'roc_auc_ovr_weighted', 'roc_auc_ovo_weighted',  
'balanced_accuracy', 'average_precision', 'neg_log_loss', 'neg_brier_score',  
'adjusted_rand_score', 'rand_score', 'homogeneity_score', 'completeness_score',  
'v_measure_score', 'mutual_info_score', 'adjusted_mutual_info_score',  
'normalized_mutual_info_score', 'fowlkes_mallows_score', 'precision',  
'precision_macro', 'precision_micro', 'precision_samples', 'precision_weighted',  
'recall', 'recall_macro', 'recall_micro', 'recall_samples', 'recall_weighted',  
'f1', 'f1_macro', 'f1_micro', 'f1_samples', 'f1_weighted', 'jaccard',  
'jaccard_macro', 'jaccard_micro', 'jaccard_samples', 'jaccard_weighted'])
```

```
[88]: # --- t-test example ---  
# are MEDV means significantly different for each value of CHAS?  
from scipy.stats import ttest_ind, ttest_1samp, ttest_rel
```

```
[86]: df.groupby(['CHAS'])[['MEDV']].mean()
```



```
[86]:          MEDV
      CHAS
0      22.093843
1      28.440000
```

```
[89]: ttest_ind(df[df['CHAS']==0]['MEDV'], df[df['CHAS']==1]['MEDV'])
```

```
[89]: TtestResult(statistic=-3.996437466090509, pvalue=7.390623170519902e-05,
df=504.0)
```

```
[ ]: # --- autoregresion ---
      # Fitting an AR(p) model via OLS yields a biased estimate
```